

# Vechile Price Prediction

## Objective:

Build a system that can predict the prices for vehicles using data on Vehicle specifications, make, etc. Explore the data to understand the features and figure out an approach.

---

## 1. Problem Statement

The task is to do EDA on dataset and build a model to predict Vechile Price based on the features provided. The challenge is to create a model that can accurately predict the outcome.

---

## 2. Data Pre-Processing

### 2.1 Data Inspection and Summary Statistics

- **Load the Dataset:** Import the dataset and review its basic structure, including column names, data types, and a few initial records.
- **Generate Summary Statistics:** Calculate key statistics (mean, median, min, max, standard deviation, etc.) to understand the primary characteristics of each column.
- Changing column names and data types

### 2.2 Data Cleaning and Feature Engineering

- **Missing Values:** Check and handle missing values if present.
- **Duplicate Values:** Check duplicate values and handle if present.

### 2.3 Outlier Treatment

- **Outlier Detection:** Identify outliers in features box plots or Z-scores and apply treatment if necessary.
- 

## 3. Exploratory Data Analysis (EDA)

### 3.1 Univariate Analysis

- **Numerical Data:** Visualize distributions with histograms and box plots.
- **Categorical Data:** Use bar charts to observe the distribution of the outcome variable.

### 3.2 Bivariate Analysis

- Create scatter plots to observe relationships between numerical features.

- Use box plots to explore how numerical features differ based on the outcome variable.

### 3.3 Multivariate Analysis

- Generate a heatmap of the correlation matrix to identify potential relationships.
- 

## 4. Model Building

### 4.1 Encoding Categorical Variables:

- Convert the Categorical columns to binary format

### 4.2 Feature Engineering

- This step involves transforming raw data into meaningful features and outcome

### 4.3 Model Training

- Split the dataset into training and testing sets.
- Scalling the data
- Use a Linear Regression to train the model on the training data.
- Model Evaluation
- Visualize the result

## 5. Advanced Modeling:

- Experiment with more complex models like RandomForest and xgboost to improve predictions.

## Import Required Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split

import warnings # to ignore warnings
warnings.filterwarnings('ignore')
```

## Load the Dataset¶

```
data = pd.read_csv('C://Users//PC//Downloads//Projects-20240722T093004Z-001//Projects//Vehicle Price Prediction//dataset.csv')
```

```
data.head() # First rows
```

```

      name \
0  2024 Jeep Wagoneer Series II
1  2024 Jeep Grand Cherokee Laredo
2  2024 GMC Yukon XL Denali
3  2023 Dodge Durango Pursuit
4  2024 RAM 3500 Laramie

      description  make
model \
0  \n  \n  Heated Leather Seats, Nav Sy...  Jeep
Wagoneer
1  All West is committed to offering every custome...  Jeep  Grand
Cherokee
2  NaN  GMC
Yukon XL
3  White Knuckle Clearcoat 2023 Dodge Durango Pur...  Dodge
Durango
4  \n  \n  2024 Ram 3500 Laramie Billet...  RAM
3500

      year  price  engine \
0  2024  74600.0  24V GDI DOHC Twin Turbo
1  2024  50170.0  OHV
2  2024  96410.0  6.2L V-8 gasoline direct injection, variable v...
3  2023  46835.0  16V MPFI OHV
4  2024  81663.0  24V DDI OHV Turbo Diesel

      cylinders  fuel  mileage  transmission  trim
body \
0  6.0  Gasoline  10.0  8-Speed Automatic  Series II
SUV
1  6.0  Gasoline  1.0  8-Speed Automatic  Laredo
SUV
2  8.0  Gasoline  0.0  Automatic  Denali
SUV
3  8.0  Gasoline  32.0  8-Speed Automatic  Pursuit
SUV
4  6.0  Diesel  10.0  6-Speed Automatic  Laramie Pickup
Truck
```

|   | doors | exterior_color          | interior_color   | drivetrain       |
|---|-------|-------------------------|------------------|------------------|
| 0 | 4.0   | White                   | Global Black     | Four-wheel Drive |
| 1 | 4.0   | Metallic                | Global Black     | Four-wheel Drive |
| 2 | 4.0   | Summit White            | Teak/Light Shale | Four-wheel Drive |
| 3 | 4.0   | White Knuckle Clearcoat | Black            | All-wheel Drive  |
| 4 | 4.0   | Silver                  | Black            | Four-wheel Drive |

data.tail() # Last records

```

                                name \
997  2024 Mercedes-Benz Sprinter 2500 Standard Roof
998          2024 Dodge Hornet Hornet R/T Plus Eawd
999          2024 Jeep Wagoneer Base
1000          2024 Nissan Murano SV Intelligent AWD
1001          2024 Chevrolet Silverado 2500 WT

```

```

                                description          make
\
997  2024 Mercedes-Benz Sprinter 2500 Cargo 144 WB ... Mercedes-Benz
998  Dealer Comments +++ Price Ends 5/31/2024 +++ A... Dodge
999  \n          \n          The ALL New Friendship CDJR ... Jeep
1000 \n          \n          CVT with Xtronic, AWD.At Tod... Nissan
1001 01u 2024 Chevrolet Silverado 2500HD Work Truck... Chevrolet

```

```

            model  year  price \
997  Sprinter 2500  2024  59037.0
998      Hornet   2024  49720.0
999   Wagoneer   2024  69085.0
1000   Murano    2024  43495.0
1001 Silverado 2500  2024  48995.0

```

```

                                engine  cylinders
fuel \
997          16V DDI DOHC Turbo Diesel          4.0
Diesel
998  4 gasoline direct injection, DOHC, Multiair va...          4.0
Gasoline

```

|          |   |     |
|----------|---|-----|
| 999      | 24V GDI DOHC Twin Turbo                           | 6.0 |
| Gasoline |   |     |
| 1000     | 6 DOHC, variable valve control, regular unlead... | 6.0 |
| Gasoline |   |     |
| 1001     | 8 gasoline direct injection, variable valve co... | 8.0 |
| Gasoline |   |     |

|        | mileage | transmission                        |
|--------|---------|-------------------------------------|
| trim \ |         |                                     |
| 997    | 10.0    | 9-Speed Automatic                   |
| Roof   |         |                                     |
| 998    | 0.0     | 6-Spd Aisin F21-250 PHEV Auto Trans |
| Eawd   |         |                                     |
| 999    | 20.0    | 8-Speed Automatic                   |
| Base   |         |                                     |
| 1000   | 6.0     | Automatic                           |
| AWD    |         | SV Intelligent                      |
| 1001   | 31.0    | Automatic                           |
| WT     |         |                                     |

|      | body         | doors | exterior_color      | interior_color \ |
|------|--------------|-------|---------------------|------------------|
| 997  | Cargo Van    | 3.0   | Arctic White        | Black            |
| 998  | SUV          | 4.0   | Acapulco Gold       | Black            |
| 999  | SUV          | 4.0   | Diamond Black       | Black            |
| 1000 | SUV          | 4.0   | Pearl White Tricoat | Graphite         |
| 1001 | Pickup Truck | 4.0   | Wheatland Yellow    | Jet Black        |

|      | drivetrain       |
|------|------------------|
| 997  | Rear-wheel Drive |
| 998  | All-wheel Drive  |
| 999  | Four-wheel Drive |
| 1000 | All-wheel Drive  |
| 1001 | Rear-wheel Drive |

## 2.Data Preprocessing¶

### 2.1 Data Inspection and Summary Statistics

```
data.shape # Rows and columns
(1002, 17)

data.ndim # Dimention of data
2
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1002 entries, 0 to 1001
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   name                   1002 non-null   object
1   description            946 non-null    object
2   make                   1002 non-null   object
3   model                  1002 non-null   object
4   year                   1002 non-null   int64
5   price                  979 non-null    float64
6   engine                 1000 non-null   object
7   cylinders              897 non-null    float64
8   fuel                   995 non-null    object
9   mileage                968 non-null    float64
10  transmission           1000 non-null   object
11  trim                   1001 non-null   object
12  body                   999 non-null    object
13  doors                  995 non-null    float64
14  exterior_color         997 non-null    object
15  interior_color         964 non-null    object
16  drivetrain             1002 non-null   object
dtypes: float64(4), int64(1), object(12)
memory usage: 133.2+ KB
```

```
data.describe() # Description of data
```

|       | year        | price         | cylinders  | mileage     | doors      |
|-------|-------------|---------------|------------|-------------|------------|
| count | 1002.000000 | 979.000000    | 897.000000 | 968.000000  | 995.000000 |
| mean  | 2023.916168 | 50202.985700  | 4.975474   | 69.033058   | 3.943719   |
| std   | 0.298109    | 18700.392062  | 1.392526   | 507.435745  | 0.274409   |
| min   | 2023.000000 | 0.000000      | 0.000000   | 0.000000    | 2.000000   |
| 25%   | 2024.000000 | 36600.000000  | 4.000000   | 4.000000    | 4.000000   |
| 50%   | 2024.000000 | 47165.000000  | 4.000000   | 8.000000    | 4.000000   |
| 75%   | 2024.000000 | 58919.500000  | 6.000000   | 13.000000   | 4.000000   |
| max   | 2025.000000 | 195895.000000 | 8.000000   | 9711.000000 | 5.000000   |

Price , cylinder,mileage columns have min value 0 which is not feasible. These records will be handled during handling of outliers

```
data.columns
```

```
Index(['name', 'description', 'make', 'model', 'year', 'price',  
      'engine',  
      'cylinders', 'fuel', 'mileage', 'transmission', 'trim', 'body',  
      'doors',  
      'exterior_color', 'interior_color', 'drivetrain'],  
      dtype='object')
```

*# Checking unique values in each column*

```
print("# unique values in name:", data['name'].nunique())  
print("# unique values in description:",  
data['description'].nunique())  
print("# unique values in make:", data['make'].nunique())  
print("# unique values in model:", data['model'].nunique())  
print("# unique values in year:", data['year'].nunique())  
print("# unique values in price:", data['price'].nunique())  
print("# unique values in engine:", data['engine'].nunique())  
print("# unique values in cylindersfuel:",  
data['cylinders'].nunique())  
print("# unique values in fuel:", data['fuel'].nunique())  
print("# unique values in mileage:", data['mileage'].nunique())  
print("# unique values in trim:", data['trim'].nunique())  
print("# unique values in body:", data['body'].nunique())  
print("# unique values in doors:", data['doors'].nunique())  
print("# unique values in exterior_color:",  
data['exterior_color'].nunique())  
print("# unique values in interior_color:",  
data['interior_color'].nunique())  
print("# unique values in drivetrain:",  
data['drivetrain'].nunique())
```

```
# unique values in name: 358  
# unique values in description: 761  
# unique values in make: 28  
# unique values in model: 153  
# unique values in year: 3  
# unique values in price: 859  
# unique values in engine: 100  
# unique values in cylindersfuel: 5  
# unique values in fuel: 7  
# unique values in mileage: 95  
# unique values in trim: 197  
# unique values in body: 8  
# unique values in doors: 4  
# unique values in exterior_color: 263  
# unique values in interior_color: 91  
# unique values in drivetrain: 4
```

## 2.2 Data Cleaning¶

Duplicate Values

```
data.isnull().sum()
```

|                |       |
|----------------|-------|
| name           | 0     |
| description    | 56    |
| make           | 0     |
| model          | 0     |
| year           | 0     |
| price          | 23    |
| engine         | 2     |
| cylinders      | 105   |
| fuel           | 7     |
| mileage        | 34    |
| transmission   | 2     |
| trim           | 1     |
| body           | 3     |
| doors          | 7     |
| exterior_color | 5     |
| interior_color | 38    |
| drivetrain     | 0     |
| dtype:         | int64 |

```
data = data.dropna() # Droppping null values
```

```
data.isnull().sum() # No Null values
```

|                |       |
|----------------|-------|
| name           | 0     |
| description    | 0     |
| make           | 0     |
| model          | 0     |
| year           | 0     |
| price          | 0     |
| engine         | 0     |
| cylinders      | 0     |
| fuel           | 0     |
| mileage        | 0     |
| transmission   | 0     |
| trim           | 0     |
| body           | 0     |
| doors          | 0     |
| exterior_color | 0     |
| interior_color | 0     |
| drivetrain     | 0     |
| dtype:         | int64 |



Duplicate values

```
data.duplicated().sum()
```

```
16
```

```
data = data.drop_duplicates() # Dropping duplicates
```

```
data.duplicated().sum() # NO duplicate values
```

```
0
```

## 2.3 Outlier Treatment¶

```
data.dtypes
```

```
name          object
description    object
make          object
model         object
year          int64
price         float64
engine        object
cylinders     float64
fuel          object
mileage       float64
transmission  object
trim          object
body         object
doors        float64
exterior_color object
interior_color object
drivetrain    object
dtype: object
```





```
fig,axis = plt.subplots(2,2,figsize=(8,10))
```

```
sns.boxplot(ax = axis[0][0], data=data['price'],palette='Set2')
```

```
sns.boxplot(ax = axis[0][1], data=data['cylinders'])
```

```
sns.boxplot(ax = axis[1][0], data=data['mileage'],palette='coolwarm')
```

```
<Axes: ylabel='mileage'>
```



*# Handling Outliers*

```
def outlier(x):
```

```
    q1 = data[x].quantile(0.25)
```

```
    q3 = data[x].quantile(0.75)
```





```
    iqr = q3-q1
```

```
lower = q1 - 1.5*iqr
upper = q3 + 1.5*iqr
median = np.median(data[x])
for i in data[x]:
    if i <lower or i>upper:
        data[x]=data[x].replace({i:median})

outlier('price')
outlier('cylinders')
outlier('mileage')

fig,axis = plt.subplots(2,2,figsize=(8,10))

sns.boxplot(ax = axis[0][0], data=data['price'],palette='Set2')
sns.boxplot(ax = axis[0][1], data=data['cylinders'])
sns.boxplot(ax = axis[1][0], data=data['mileage'],palette='coolwarm')
<Axes: ylabel='mileage'>
```



### 3. EDA

#### 3.1 Univariate Analysis


Visualize individual variables to understand their distribution (e.g., histograms for numerical data, bar charts for categorical data).

#### 3.2 Bivariate and Multivariate Analysis

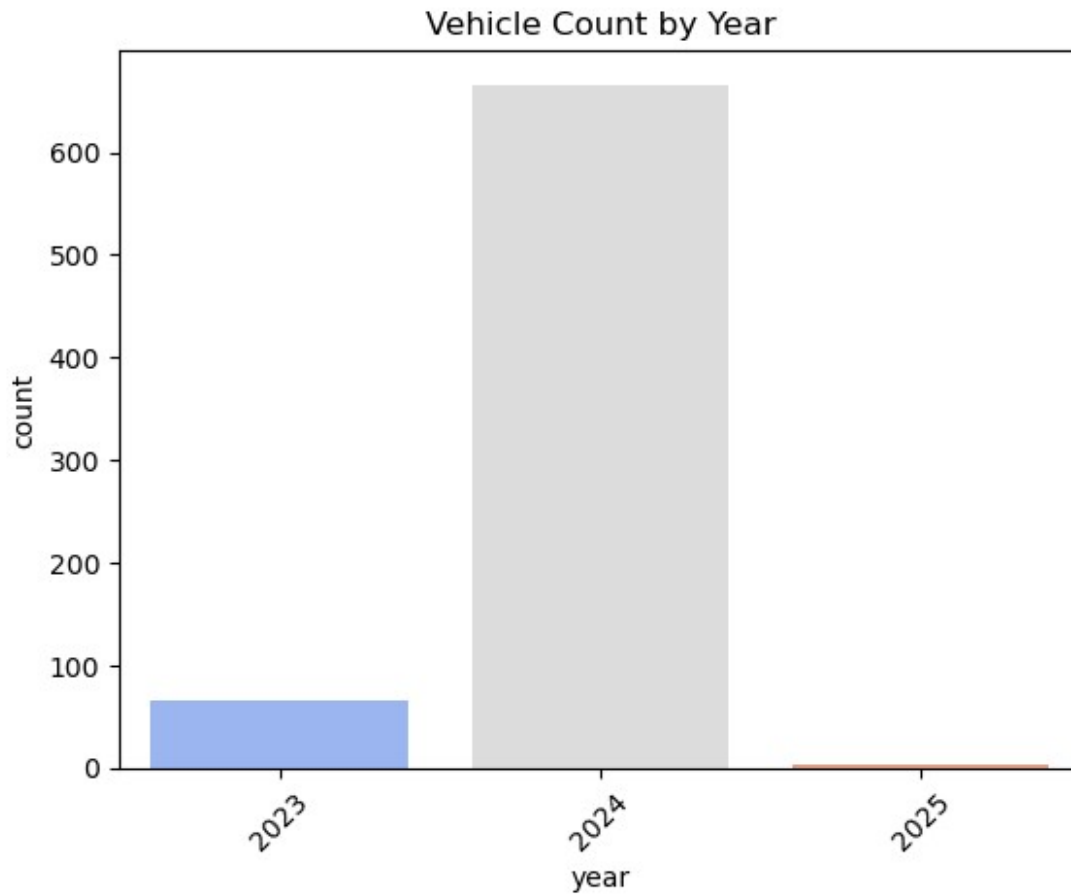
Explore relationships between variables by visualizing pairs of variables or groups of variables (e.g., scatter plots, heatmaps).

#### 3.1 Univariate Analysis

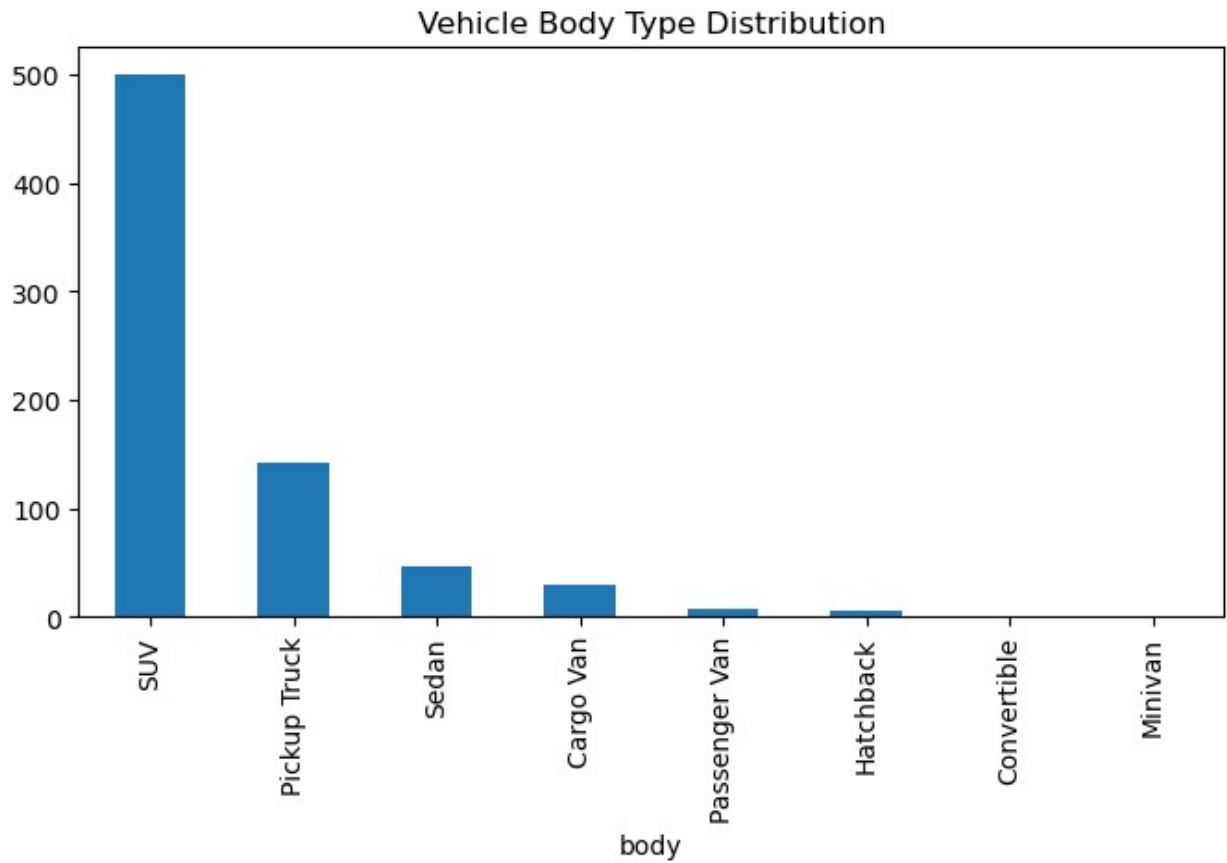
```
plt.figure(figsize=(10, 5))
sns.histplot(data['price'], bins=50, kde=True)
plt.title('Distribution of Vehicle Prices')
Text(0.5, 1.0, 'Distribution of Vehicle Prices')
```



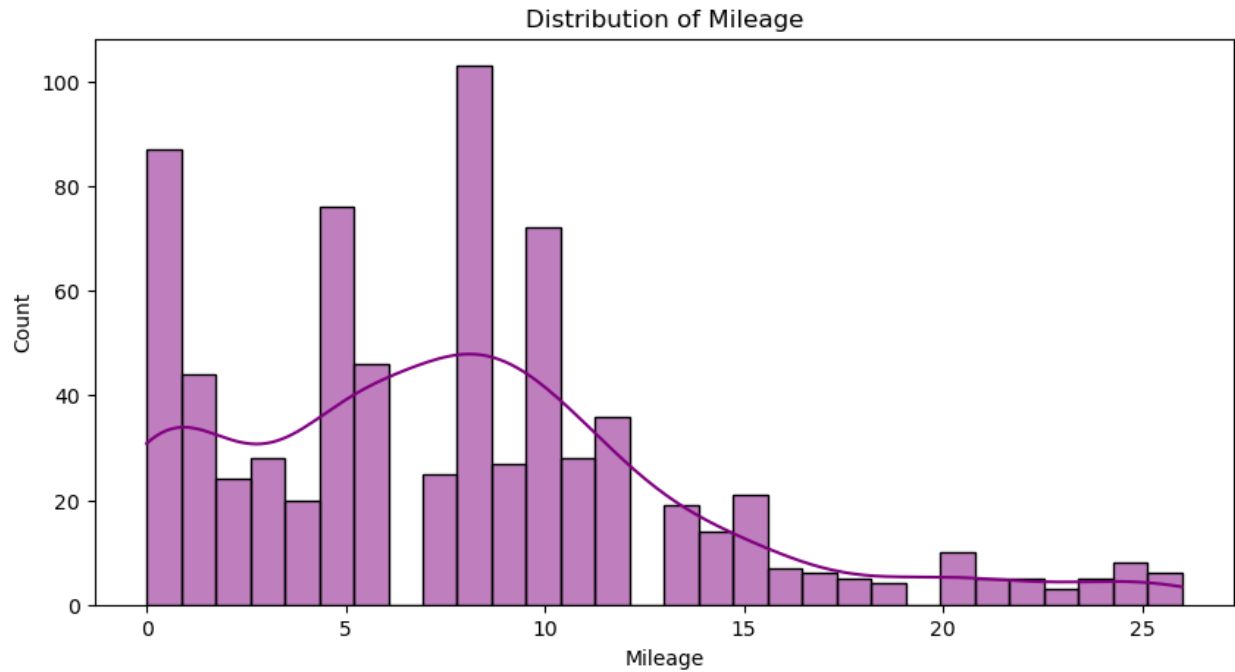
```
sns.countplot(data=data, x='year',palette='coolwarm')
plt.xticks(rotation=45)
plt.title('Vehicle Count by Year')
Text(0.5, 1.0, 'Vehicle Count by Year')
```



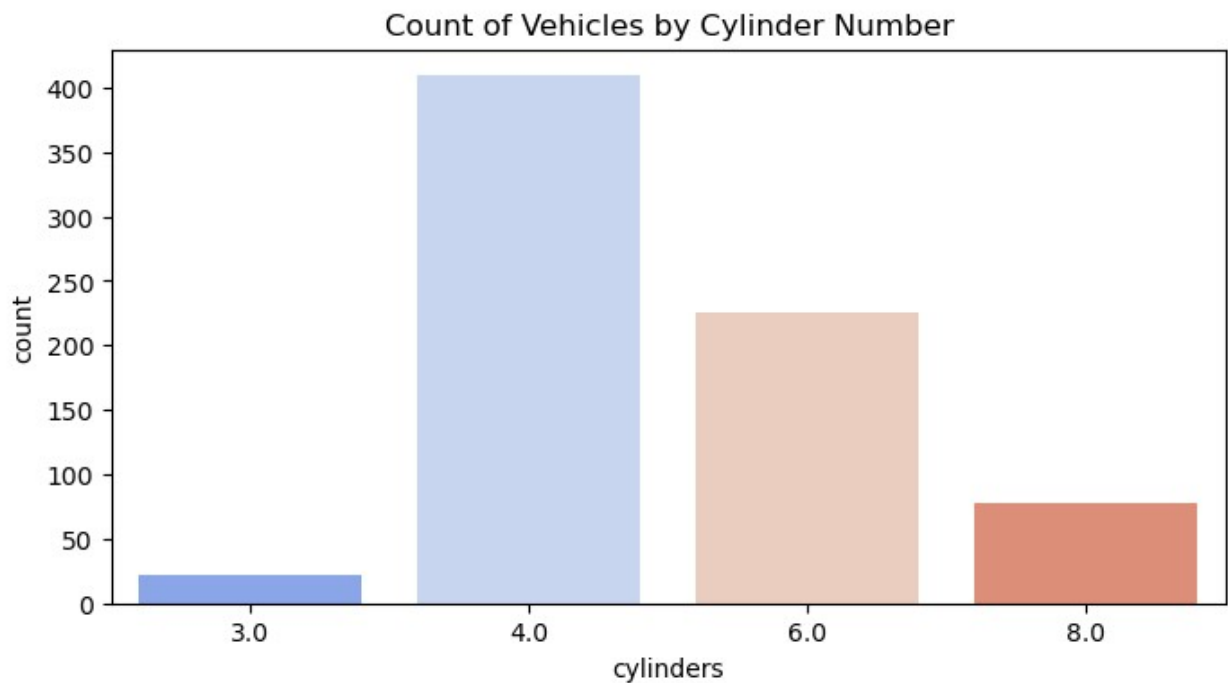
```
data['body'].value_counts().plot(kind='bar', figsize=(8, 4))
plt.title('Vehicle Body Type Distribution')
Text(0.5, 1.0, 'Vehicle Body Type Distribution')
```



```
plt.figure(figsize=(10, 5))
sns.histplot(data['mileage'], bins=30, kde=True, color='purple')
plt.title('Distribution of Mileage')
plt.xlabel('Mileage')
Text(0.5, 0, 'Mileage')
```




```
plt.figure(figsize=(8, 4))
sns.countplot(data=data, x='cylinders', palette='coolwarm')
plt.title('Count of Vehicles by Cylinder Number')
Text(0.5, 1.0, 'Count of Vehicles by Cylinder Number')
```






```
sns.countplot(data=data, x='doors', palette='viridis')
plt.title('Distribution of Number of Doors')
Text(0.5, 1.0, 'Distribution of Number of Doors')
```




### 3.2 Bivariate analysis

```
sns.boxplot(data=data, x='year', y='price')
plt.title('Price by Year')
Text(0.5, 1.0, 'Price by Year')
```



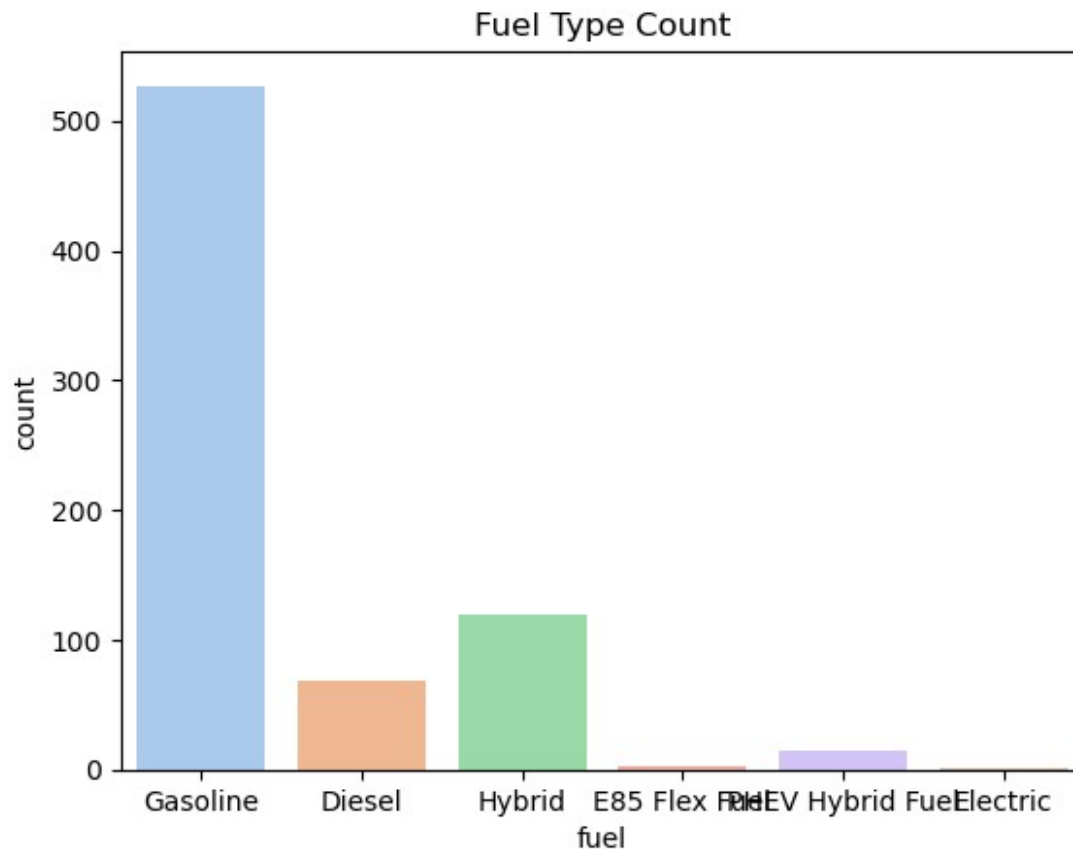
```
plt.figure(figsize=(12, 6))
sns.boxplot(data=data, x='engine', y='price', palette='Set2')
plt.xticks(rotation=90)
plt.title('Price by Engine Type')
Text(0.5, 1.0, 'Price by Engine Type')
```

Price by Engine Type

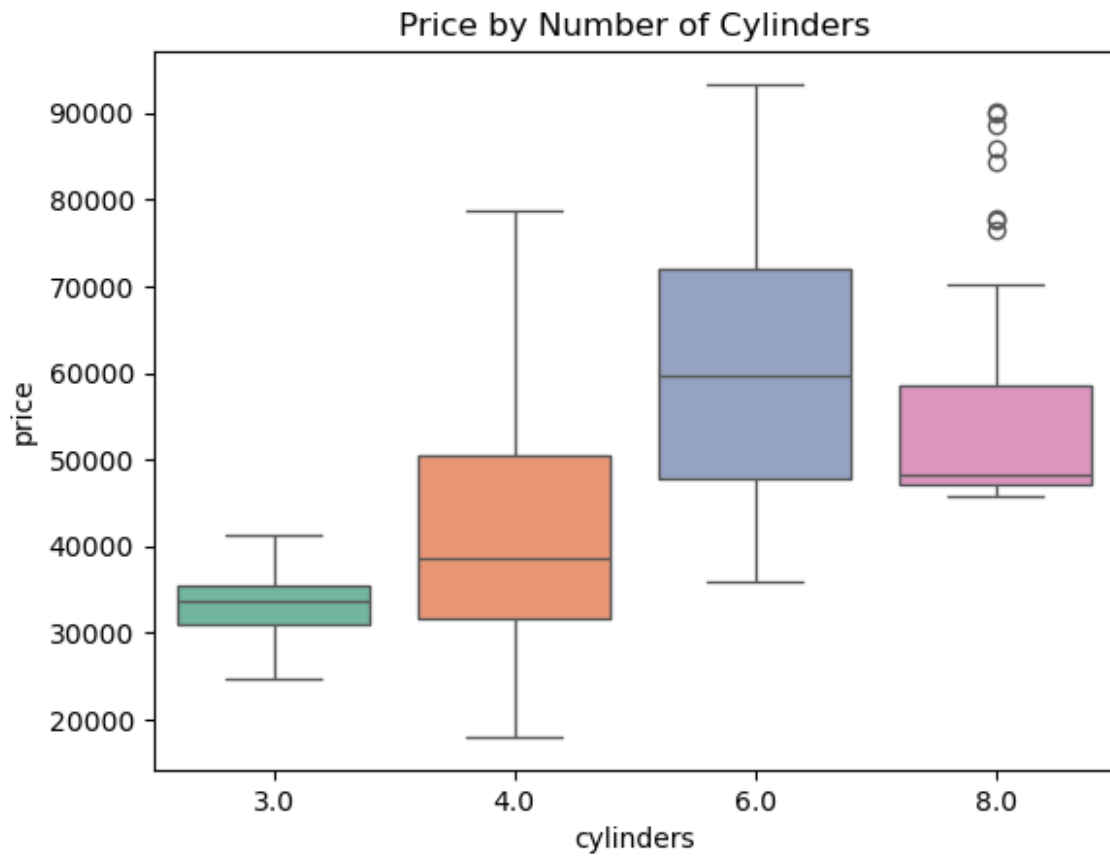


```
sns.countplot(data=data, x='fuel', palette='pastel')
plt.title('Fuel Type Count')
```


```
Text(0.5, 1.0, 'Fuel Type Count')
```



```
sns.boxplot(data=data, x='cylinders', y='price', palette='Set2')  
plt.title('Price by Number of Cylinders')  
Text(0.5, 1.0, 'Price by Number of Cylinders')
```



```
sns.boxplot(data=data, x='transmission', y='price',palette='coolwarm')
plt.xticks(rotation=90)
plt.title('Price by Transmission Type')
Text(0.5, 1.0, 'Price by Transmission Type')
```

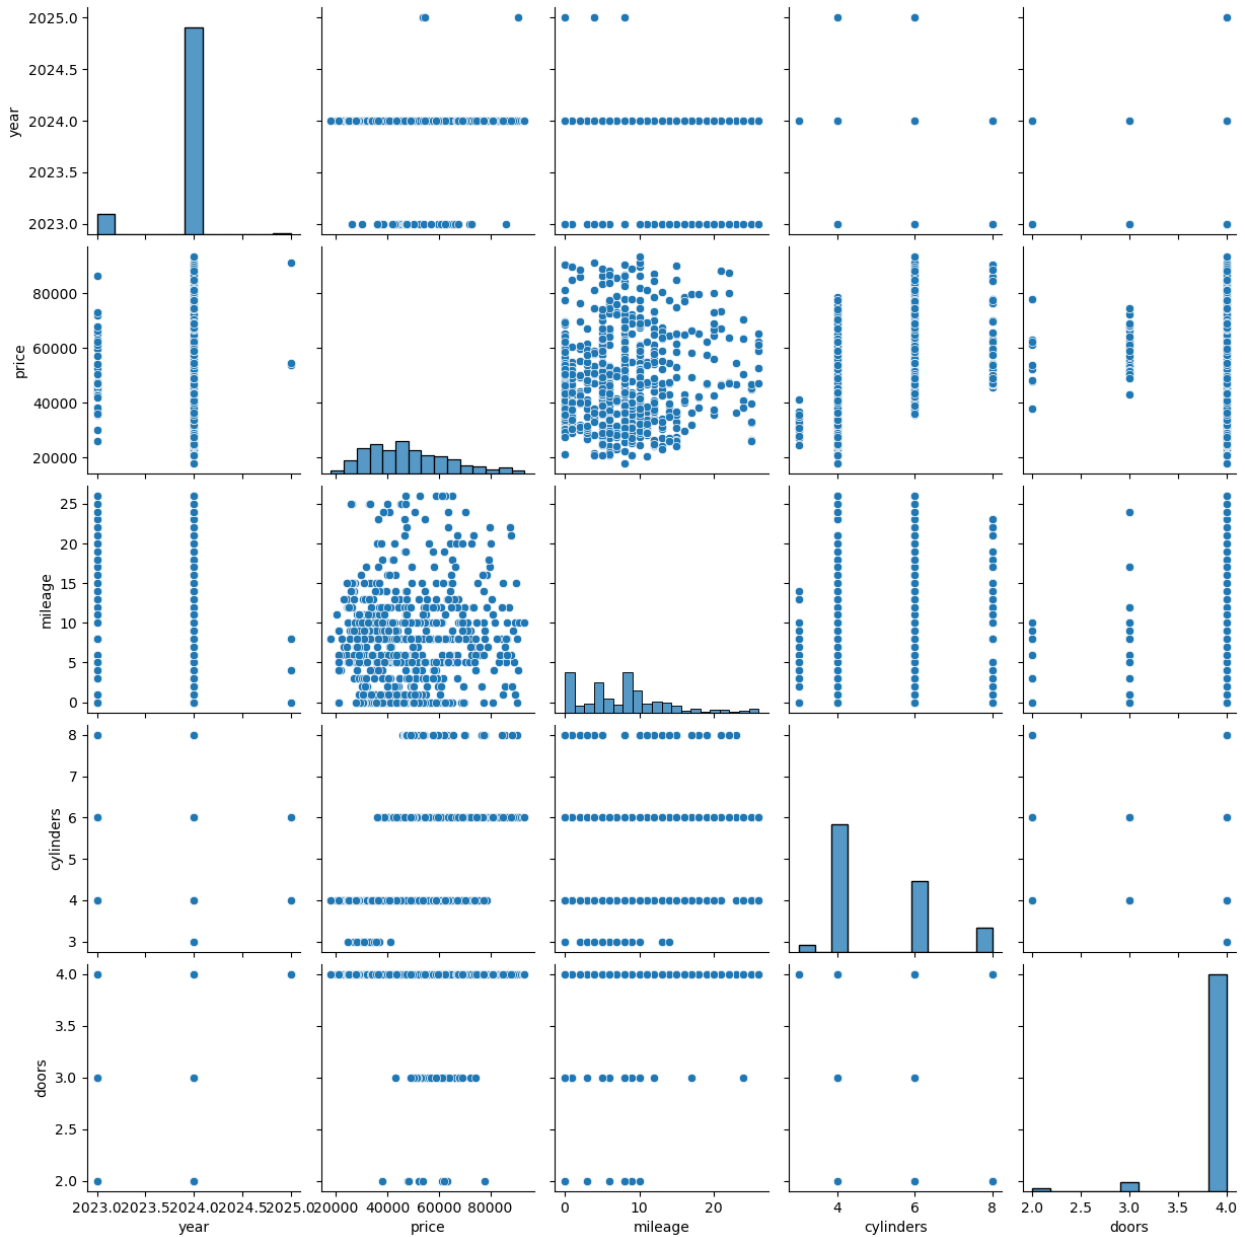


```
num_features = ['year', 'price', 'mileage', 'cylinders', 'doors']

# Pairplot for correlations
```

```
sns.pairplot(data[num_features])
```


```
<seaborn.axisgrid.PairGrid at 0x1a0d01434a0>
```



### 3.3 Multivariate Analysis

```
plt.figure(figsize=(10, 6))
sns.heatmap(data.corr(numeric_only=True), annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
```


```
Text(0.5, 1.0, 'Correlation Matrix')
```



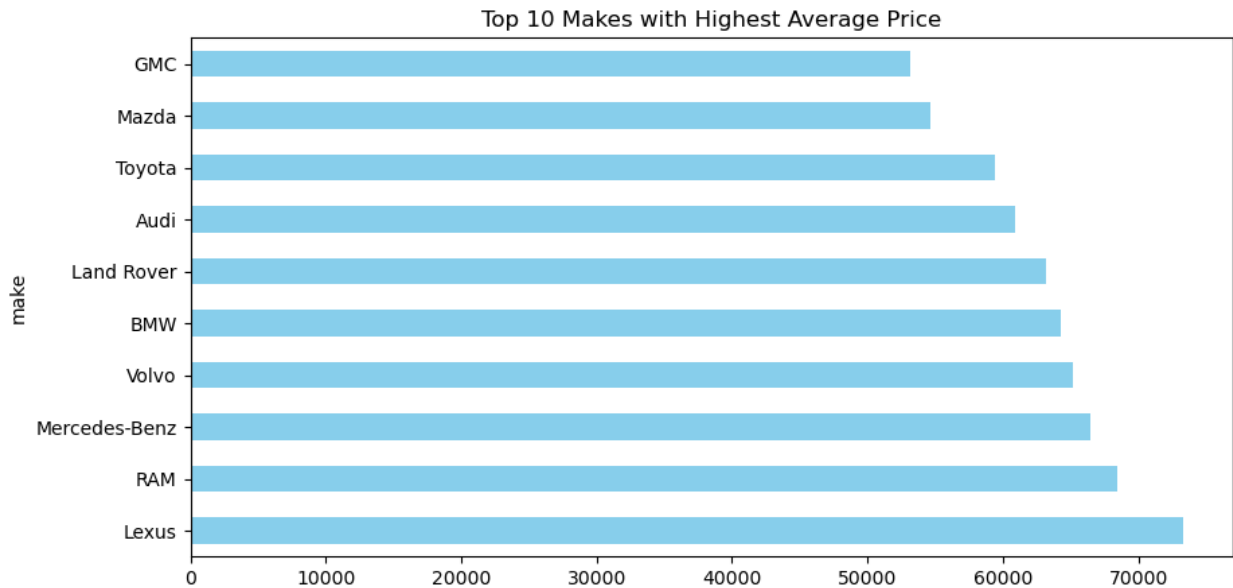
```
plt.figure(figsize=(10, 6))
sns.scatterplot(data=data, x='mileage', y='price', hue='drivetrain')
plt.title('Price vs Mileage (by Drivetrain)')
```

```
Text(0.5, 1.0, 'Price vs Mileage (by Drivetrain)')
```





```
avg_price_by_make = data.groupby('make')
['price'].mean().sort_values(ascending=False).head(10)
avg_price_by_make.plot(kind='barh', figsize=(10, 5), color='skyblue')
plt.title('Top 10 Makes with Highest Average Price')
Text(0.5, 1.0, 'Top 10 Makes with Highest Average Price')
```



Grouped rare categories in model, engine, and trim features into 'Other' to reduce cardinality and improve model generalization.

```
model_counts = data['model'].value_counts()
common_models = model_counts[model_counts > 5].index
data['model'] = data['model'].apply(lambda x: x if x in common_models
else 'Other')

engine_counts = data['engine'].value_counts()
common_engine = engine_counts[engine_counts > 5].index

# Replace engine models with 'Other'
data['engine'] = data['engine'].apply(lambda x: x if x in
common_engine else 'Other')

trim_counts = data['trim'].value_counts()
common_trim = trim_counts[trim_counts > 3].index

# Replace rare Trim with 'Other'
data['trim'] = data['trim'].apply(lambda x: x if x in common_trim else
'Other')
```

## 4. Model Building

### 4.1 Encoding Categorical columns

```
data =  
pd.get_dummies(data, columns=['make', 'fuel', 'body', 'doors', 'drivetrain',  
                             'model', 'engine', 'trim'], drop_first=True)  
data.head()
```

```
      name \  
0  2024 Jeep Wagoneer Series II  
1  2024 Jeep Grand Cherokee Laredo  
3    2023 Dodge Durango Pursuit  
4    2024 RAM 3500 Laramie  
5    2024 Nissan Murano Platinum
```

```
      description  year  price \  
0  \n      \n      Heated Leather Seats, Nav Sy...  2024  74600.0  
1  Al West is committed to offering every custome...  2024  50170.0  
3  White Knuckle Clearcoat 2023 Dodge Durango Pur...  2023  46835.0  
4  \n      \n      2024 Ram 3500 Laramie Billet...  2024  81663.0  
5  \n      \n      Boasts 28 Highway MPG and 20...  2024  46000.0
```

```
      cylinders  mileage  transmission  exterior_color \  
0           6.0    10.0  8-Speed Automatic           White  
1           6.0     1.0  8-Speed Automatic           Metallic  
3           8.0     8.0  8-Speed Automatic  White Knuckle Clearcoat  
4           6.0    10.0  6-Speed Automatic           Silver  
5           6.0     8.0    Automatic CVT           White
```

```
      interior_color  make_BMW  ...  trim_Sahara  trim_Series II  
trim_Series III \  
0  Global Black      False  ...      False      True  
False  
1  Global Black      False  ...      False      False  
False  
3           Black      False  ...      False      False  
False  
4           Black      False  ...      False      False  
False  
5           Gray      False  ...      False      False  
False
```

```
      trim_Sport  trim_Sport S  trim_Tradesman \  
0           False           False           False  
1           False           False           False  
3           False           False           False
```

|   |       |       |       |
|---|-------|-------|-------|
| 4 | False | False | False |
| 5 | False | False | False |

|   |  |                |            |
|---|--|----------------|------------|
|   | trim_Tradesman Crew Cab 4x4 8&#39; Box | trim_Trailhawk | trim_XLT \ |
| 0 | False                                  | False          | False      |
| 1 | False                                  | False          | False      |
| 3 | False                                  | False          | False      |
| 4 | False                                  | False          | False      |
| 5 | False                                  | False          | False      |

|   |                |
|---|----------------|
|   | trim_xDrive40i |
| 0 | False          |
| 1 | False          |
| 3 | False          |
| 4 | False          |
| 5 | False          |

[5 rows x 164 columns]

## 4.2 Feature Engineering

```
X =
data.drop(columns=['name', 'description', 'price', 'exterior_color', 'interior_color', 'transmission'])
y = data['price']
```

## 4.3 Model Training

```
# Split the data into training and testing

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Standarize the data

scaler = StandardScaler()
X_train_normalized = scaler.fit_transform(X_train)
X_test_normalized = scaler.transform(X_test)
```

```
# Using Linear Regression to predict profit
```

```
model = LinearRegression()
```

```
# Training the model
```

```
model.fit(X_train_normalized, y_train)
```

```
# Make Prediction
```

```
y_pred = model.predict(X_test_normalized)
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
rmse = np.sqrt(mse)
```

```
r2 = r2_score(y_test, y_pred)
```

```
print(f"Root Squared Error: {rmse}")
```

```
print(f"R-squared: {r2}")
```

```
Root Squared Error: 5720.5809941552125
```

```
R-squared: 0.8814003113749823
```

```
plt.figure(figsize=(8, 6))
```

```
plt.scatter(y_test, y_pred)
```

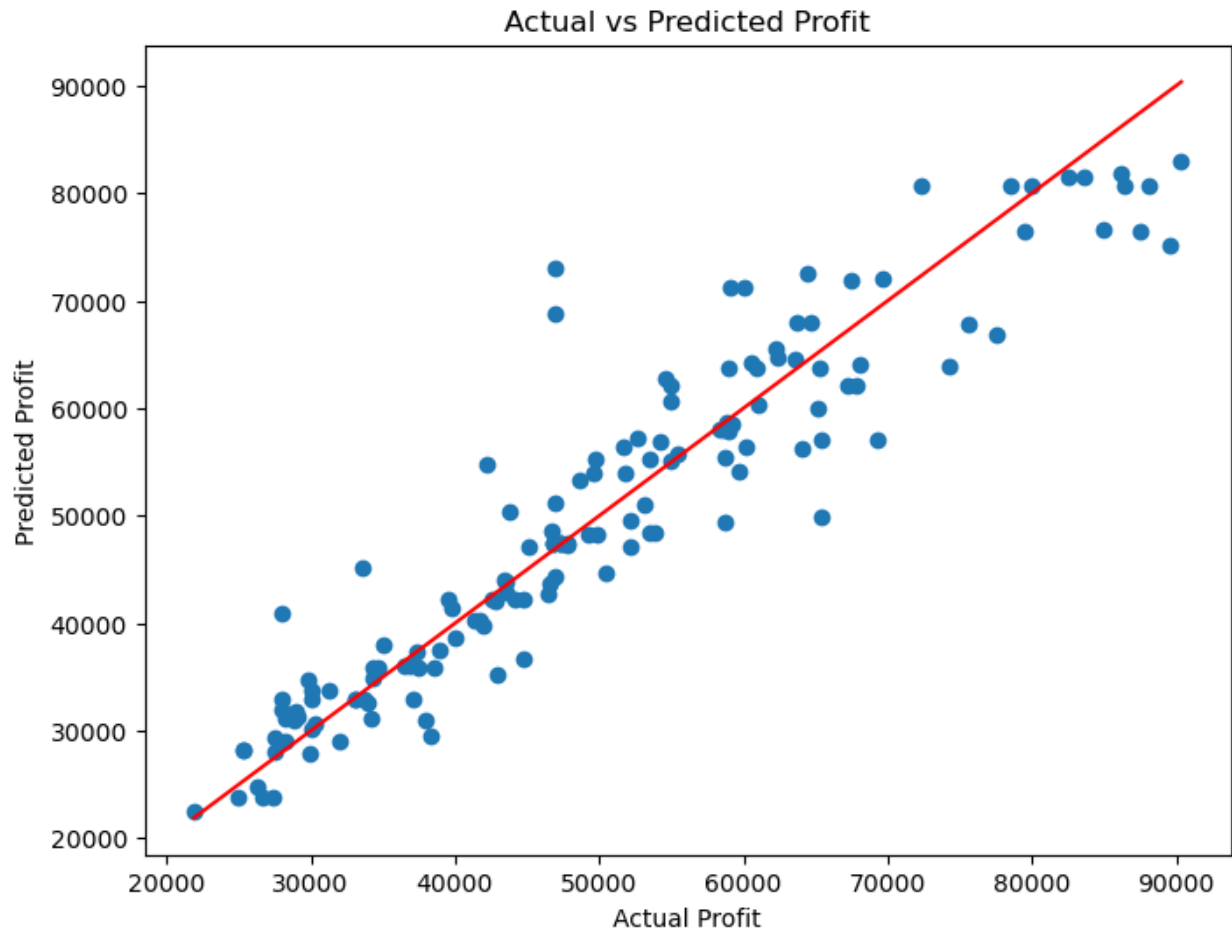
```
plt.plot([min(y_test), max(y_test)], [min(y_test),  
max(y_test)], color='red')
```

```
plt.title('Actual vs Predicted Profit')
```

```
plt.xlabel('Actual Profit')
```

```
plt.ylabel('Predicted Profit')
```

```
plt.show()
```



## 5. Advance Model

*# Random Forrest*

```
from sklearn.ensemble import RandomForestRegressor
rf = RandomForestRegressor(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
y_pred_log = rf.predict(X_test)
```

```
mse = mean_squared_error(y_test, y_pred_log)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred_log)
```

```
print(f"Root Squared Error: {rmse}")
print(f"R-squared: {r2}")
```

```
Root Squared Error: 5692.849758133807
R-squared: 0.8825473781654893
```

```
# XGBoost
```

```
from xgboost import XGBRegressor
xgb = XGBRegressor(n_estimators=100, learning_rate=0.1,
random_state=42)
xgb.fit(X_train, y_train)
y_pred_xgb = xgb.predict(X_test)
```

```
mse = mean_squared_error(y_test, y_pred_xgb)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred_xgb)
```

```
print(f"Root Squared Error: {rmse}")
print(f"R-squared: {r2}")
```

```
Root Squared Error: 5839.966089939761
R-squared: 0.8763984486309568
```

## Conclusion

To identify the best-performing algorithm for vehicle price prediction, we evaluated three models: Linear Regression, Random Forest, and XGBoost. Among

them, the Random Forest model delivered the most accurate results, achieving the lowest Root Mean Squared Error (RMSE) of 5692.85 and the highest R-

squared value of 0.8825. This indicates that the model captures approximately 88.25% of the variance in vehicle prices, making it the most reliable

choice. While Linear Regression and XGBoost also performed well, Random Forest showed a slightly better balance between prediction accuracy and variance explanation.